

Green Resource Allocation Algorithms for Publish/Subscribe Systems

Distributed Systems Course Paper Presentation

Artem Tsikiridis

December 14, 2015

`cs.grad.aueb`

Outline

Background

Overview/Prerequisites

Phase One

Phase Two

Phase Three

Experiment

Conclusions

Background

Green/Energy Saving Algorithms

- Green → sustainable energy practices
- Green IT initiatives are a priority

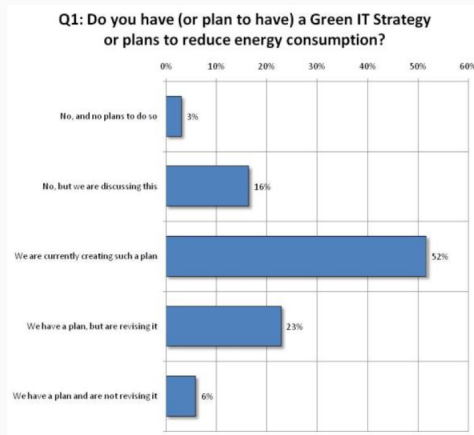


Figure 1: source: Symantec Green IT Report, 2009

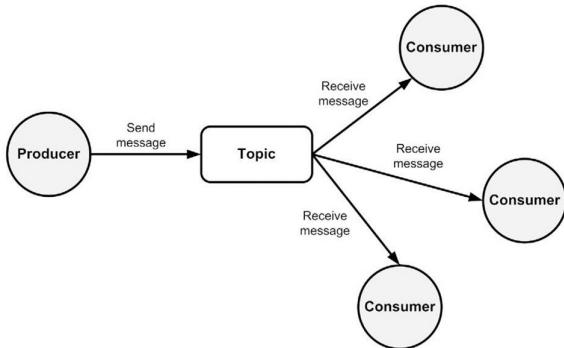
Efficient use of resources leads to **lower** IT operational costs.

Message Broker Pattern

- Module translating messages from messaging protocol of the sender to messaging protocol of the receiver
- Widely used in computer networks
- Minimizes mutual awareness of applications → effective decoupling
- Apache ActiveMQ, Apache Kafka, Celery



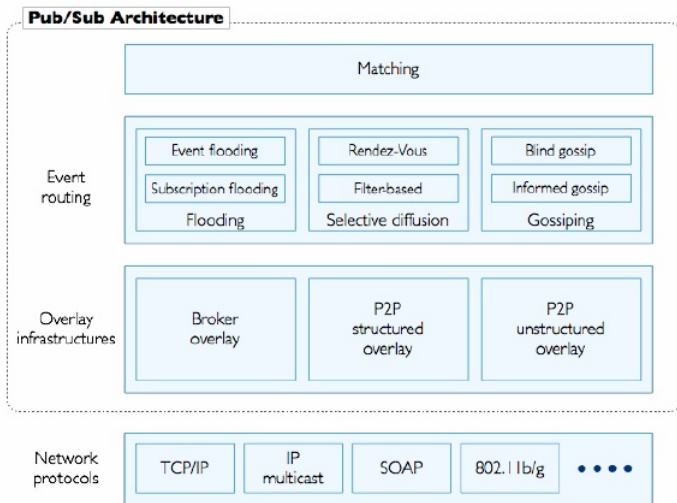
Publish/Subscribe Systems



- Publishers (senders) put messages to classes
- Subscribers (receivers) express interest in classes of messages
- Broker-based systems: publishers post messages to message broker, subscribers register subscriptions with that broker
- Used by Green IT implementors, provides network scalability

Generic Publish/Subscribe Systems

■ A generic architecture of a publish/subscribe system:



From "Distributed Event Routing in Publish/Subscribe Communication Systems: a survey"
R. Baldoni, L. Querzoni, S. Takoma, A. Virgillito midlab tech.rep. 2007, to appear (springer)

Overview/Prerequisites

Main Goal: **minimize** number of brokers, while **maximizing** resource utilization of allocated brokers

- To minimize number of brokers → minimize number of messages
- By minimizing number of brokers → minimize size of network
- As a result, publication hop count is improved (less complexity to place a publication)

How? Design and implement a 3-phase scheme to **reconfigure** the publish/subscribe system.

Phases of 3-phase scheme

- Phase One: Gathering of performance and workload information from the network using bit-vectors.
- Phase Two: Allocation of subscriptions to brokers using the info from Phase One
- Phase Three: Recursively construct broker overlay with already allocated subscriptions
- After 3-phase scheme: Strategically relocate publishers to new broker overlay

Solving an optimization problem which is proven to be \mathcal{NP} -complete.

Required components for all phases

- **CROC:** Coordinator for Reconfiguring the Overlay and Clients
 - External pub/sub client
 - Connects to any broker to collect info
 - Executes phases 2 and 3
 - Orchestrates reconfiguration
- **CBC:** CROC Back-end Component
 - Integrated into each broker
 - Responds to commands sent by CROC (information requests etc.)

Phase One

Phase One: Information Gathering

- Information gathering protocol implemented using publish/subscribe.
- CROC connects to first broker → sends *Broker Information Request* message
- Broker broadcasts it to all neighbors
- Broker replies to CROC with a *Broker Information Answer* when:
 - No neighbors to forward the BIR
 - Has recieved all BIA from other neighbors

Phase One: BIA

A Broker Information Answer (BIA) consists of:


- **URL:** Useful to reassign subscribers in phases 2 and 3
- **Matching delay function:** Linear function enabling CROC to predict input load of broker
- **Total Output Bandwidth:** CROC uses it to predict output load of the broker
- **Local subscriptions and profiles:** CROC may relocate subscriptions based on that
- **Local publishers and profiles:** CROC uses this info to predict load imposed by each subscription

Once CROC gets all BIA it executes Phase 2 (reassignment of subscriptions to brokers) and Phase 3(reconfiguration of broker overlay).

Phase One: Subscription and Publisher Profiles

Subscription Profile: Captures publications sinked by subscription.

- Allows CROC to accurately estimate load requirements.
- Generated by broker's CBC



Phase Two

Phase Two: Subscription allocation algorithms

Three proposed algorithms:

- Fastest Broker First (FSF)
- Bin Packing
- Clustering with Resource Awareness and Minimization (CRAM)

subscription pool: All subscriptions reported in BIA in Phase 1.

broker pool: All brokers that sent a BIA message back to CROC

Algorithmic Input: subscription pool, broker pool

Algorithmic Output: set of non-connected brokers (possibly) with allocated subscriptions

Phase Two: Fastest Broker First

Data: subscription pool, broker pool

Result: set of non-connected brokers (possibly) with subscriptions

Sort (desc) brokers by total available bandwidth

while *Subscription pool not empty* **do**

 s = Randomly pick a subscription

 Remove s from pool

 Assign s to first broker in list that **"can handle it"**

if *s cannot be allocated to any broker* **then**

 return broker pool

end

end

return broker pool

- **broker "can handle" subscription:** Remaining output bandwidth > 0 and incoming publication rate \leq max matching rate (BIA Matching delay)
- **Complexity:** $O(S)$, S number of subscriptions

Phase Two: Bin Packing

Data: subscription pool, broker pool

Result: set of non-connected brokers (possibly) with subscriptions

Sort (desc) brokers by total available bandwidth

Sort (desc) subscriptions by bandwidth requirement

while *Subscription pool not empty* **do**

 s = Pick first subscription

 Remove s from pool

 Assign s to first broker in list that **"can handle it"**

if *s cannot be allocated to any broker* **then**

 return broker pool

end

end

return broker pool

- **Complexity:** $O(S \log(S))$, S number of subscriptions
- FBF and Bin Packing \rightarrow *sorting* algorithms

Phase Two: CRAM Metrics

- CRAM is significantly different!
- Relies on the following subscription **closeness** metrics:
 - **INTERSECT**: $|S_1 \cap S_2|$
 - **XOR (inversed)**: $|S_1 \oplus S_2|^{-1}$
 - **IOS**: $\frac{|S_1 \cap S_2|^2}{|S_1| + |S_2|}$
 - **IOU**: $\frac{|S_1 \cap S_2|^2}{|S_1 \cup S_2|}$

where S_1 and S_2 are two subscriptions.

Phase Two: CRAM Algorithm

Data: subscription pool, broker pool

Result: set of non-connected brokers (possibly) with subscriptions

Run Bin Packing algorithm on Data

while *True* **do**

 s1, s2 = Pick two closest subscriptions (using metrics)

if *s1 or s2 empty* **then**

 | return broker pool

end

 s' = s1 or s2 (bitwise)

 Remove s1 and s2 from pool

 Allocate s' using Bin Packing algorithm

if *Allocation of s' fails* **then**

 | s1, s2 = Revert s1 or s2 (bitwise)

 | Note s1, and s2 should not be clustered again


end

end

- Complexity: $O(S^3 \log(S))$
- Optimizations proposed...

Phase Two: CRAM Optimizations

1. **Grouping of Equal Subscriptions:** Consider subscriptions of equal bit vectors identical. Reduce S .
2. **Search Pruning:** Reduce search space for closeness.
 $O(S^3 \log(S)) \rightarrow O(S^2 \log(S))$
3. **One-to-Many Clustering:** Cluster S_1 with shaded subscriptions before clustering with S_2 (even though closeness is worse).




Phase Three

Phase Three: Recursive Overlay Construction


- Design a tree overlay of connected brokers.
- Use **CROC** to apply reconfiguration.
- Finally, use **GRAPE** algorithm to relocate publishers to the final overlay (out of scope of this paper).
- Three proposed optimizations for overlay construction.

Phase Three: Optimization One




(a) Opt. 1: Deallocate B5 since it is a pure forwarding broker

Phase Three: Optimization Two



(b) Opt. 2: B3 has enough capacity to serve B2's subscriptions directly along with B1

Phase Three: Optimization Three



(c) Opt. 3: Replace under-utilized high resource brokers with low resource brokers

Experiment

- Comparison of proposed framework with related approaches and baseline algorithms
- Usage of PADRES, an open-source distributed content-based publish/subscribe system implemented at UToronto

Conclusions

Conclusions

- The approach works on any implementation or variation of a pub/sub system.
- Advantages:
 1. CRAM reduces the average broker message rate by up to 92% comparing with prior work.
 2. CRAM reduces the number of allocated brokers by up to 91%, no publication hop.
 3. IOU metric heavily reduces computation time of CRAM
- Disadvantages
 1. CRAM is computationally heavier than all other approaches
 2. CRAM may set high delivery delays
- Possible Remedies
 1. Adjustments to bit vector length (reduce computation time)
 2. Strategically relocate publishers to reduce delivery delays

Thank you! Questions?